

Sekhar__Mekala__HW5

Sekhar Mekala

Tuesday, November 01, 2016

Problem 1

a)

The transition matrix is given below:

```
p <- matrix(c(.9,.05,.03,.02,0,.85,.09,.06,0,0,.9,.1,1,0,0,0),byrow=TRUE,nrow=4)
print(p)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]  0.9 0.05 0.03 0.02
## [2,]  0.0 0.85 0.09 0.06
## [3,]  0.0 0.00 0.90 0.10
## [4,]  1.0 0.00 0.00 0.00
```

b)

Given that the new state always starts in low state. Therefore the probabilities of various states after 3 weeks are given below:

```
library(expm)
```

```
## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:base':
##
##      crossprod, tcrossprod
##
## Attaching package: 'expm'
##
## The following object is masked from 'package:Matrix':
##
##      expm
```

```
#Initial state
x <- matrix(c(1,0,0,0),byrow=TRUE,nrow=1)
x %*% (p %^% 3)
```

```
##      [,1]      [,2]      [,3]      [,4]
## [1,] 0.771 0.115875 0.085425 0.0277
```

Therefore, the probability that the machine will be in failed state after 3 weeks is 0.0277

c)

Let us find the probabilities of failures in the first, second and third weeks:

First week's probability of failure is obtained as follows:

```
x %*% (p %^1)[,4]
```

```
##      [,1]
```

```
## [1,] 0.02
```

Second week's probability of failure is obtained as follows:

```
x %*% (p %^2)[,4]
```

```
##      [,1]
```

```
## [1,] 0.024
```

Third week's probability of failure is obtained as follows:

```
x %*% (p %^3)[,4]
```

```
##      [,1]
```

```
## [1,] 0.0277
```

Let us assume the following events:

- $E1$ that the machine fails in the first week. Its probability = $P(E1) = 0.02$
- $E2$ that the machine fails in the second week. Its probability = $P(E2) = 0.024$
- $E3$ that the machine fails in the third week. Its probability = $P(E3) = 0.0277$

Therefore the probability of $E1$ or $E2$ or $E3$ is obtained as:

$$P(E_1 \cup E_2 \cup E_3) = P(E_1) + P(E_2) + P(E_3) - P(E_1 \cap E_2) - P(E_2 \cap E_3) - P(E_1 \cap E_3) + P(E_1 \cap E_2 \cap E_3)$$

Since $E1$, $E2$ and $E3$ are independent,

$$P(E_1 \cap E_2) = (0.02)(0.024) = 0.00048$$

$$P(E_2 \cap E_3) = (0.024)(0.0277) = 0.0006648$$

$$P(E_1 \cap E_3) = (0.02)(0.0277) = 0.000554$$

$$P(E_1 \cap E_2 \cap E_3) = (0.02)(0.024)(0.0277) = 0.000013296$$

Hence, the probability of at least one failure in first 3 weeks is 0.0700145

d)

Let us find the steady state probability can be found by solving the following equation, where a, b, c and d represent the low, medium, high and failed states probabilities respectively in steady state:

$$\begin{bmatrix} a & b & c & d \end{bmatrix} \begin{bmatrix} 0.9 & 0.05 & 0.03 & 0.02 \\ 0.0 & 0.85 & 0.09 & 0.06 \\ 0.0 & 0.00 & 0.90 & 0.10 \\ 1.0 & 0.00 & 0.00 & 0.00 \end{bmatrix} = \begin{bmatrix} a & b & c & d \end{bmatrix}$$

After solving the above equation, we obtained

$$a = p(\text{low}) = 0.4918033$$

$$b = p(\text{medium}) = 0.1639344$$

$$c = p(\text{high}) = 0.295082$$

$$d = p(\text{failed}) = 0.04918033$$

Hence the average number of weeks for the first failure = $1/0.04918033 = 20.3333325$ weeks

e)

On an average the number of weeks the machine works in 1 year = $52 \times 0.4918033 = 25.5737716$ weeks

f)

Average profit in long run is:

$$\begin{bmatrix} 0.4918033 & 0.1639344 & 0.295082 & 0.04918033 \end{bmatrix} \cdot \begin{bmatrix} 1000 \\ 500 \\ 400 \\ -700 \end{bmatrix}$$

```
matrix(c(0.4918033, 0.1639344, 0.295082, 0.04918033),nrow=1,byrow=TRUE) %*%
matrix(c(1000, 500, 400, -700),nrow=4,byrow=TRUE)
```

```
##           [,1]
## [1,] 657.3771
```

Hence, a profit of \$657.3771 is obtained per week.

g)

When the machine is repaired soon after it is in the high state, then the transition matrix will be:

```
p <- matrix(c(.9,.05,.03,.02,0,.85,.09,.06,1,0,0,0,1,0,0,0),byrow=TRUE,nrow=4)
p
```

```
##      [,1] [,2] [,3] [,4]
## [1,]  0.9 0.05 0.03 0.02
## [2,]  0.0 0.85 0.09 0.06
## [3,]  1.0 0.00 0.00 0.00
## [4,]  1.0 0.00 0.00 0.00
```

Let us find the steady state probability for the new transition matrix. It can be found by solving the following equation, where a, b, c and d represent the low, medium, high and failed states probabilities respectively in steady state:

$$\begin{bmatrix} a & b & c & d \end{bmatrix} \begin{bmatrix} 0.9 & 0.05 & 0.03 & 0.02 \\ 0.0 & 0.85 & 0.09 & 0.06 \\ 1.0 & 0 & 0 & 0 \\ 1.0 & 0.00 & 0.00 & 0.00 \end{bmatrix} = \begin{bmatrix} a & b & c & d \end{bmatrix}$$

After solving the above equation, we obtained

$$a = p(\text{low}) = 0.6976744$$

$$b = p(\text{medium}) = 0.2325581$$

$$c = p(\text{high}) = 0.04186047$$

$$d = p(\text{failed}) = 0.02790698$$

Hence, the average profit per week will be:

$$\begin{bmatrix} 0.6976744 & 0.2325581 & 0.04186047 & 0.02790698 \end{bmatrix} \cdot \begin{bmatrix} 1000 \\ 500 \\ 400 \\ -700 \end{bmatrix}$$

```
matrix(c(0.6976744, 0.2325581, 0.04186047, 0.02790698),nrow=1,byrow=TRUE) %*%
matrix(c(1000, 500, -600, -700),nrow=4,byrow=TRUE)
```

```
##      [,1]
## [1,] 769.3023
```

Hence, a profit of \$769.3023 is obtained per week. Since the profit is more in this case, it is suggested to repair the machine soon after it reaches the high state.

Problem-2

The following R code uses Metropolis-Hastings algorithm to estimate the θ value:

```
#Initialize the observed counts
count <- c(125,18,20,34)

#Create a probability function
prob <- function(y, count)
{
  #Computes the target density
  if(y < 0 || y >= 1)
```

```

    return(0)
    return((2+y)^count[1]*(1-y)^(count[2]+count[3])*y^count[4])
}

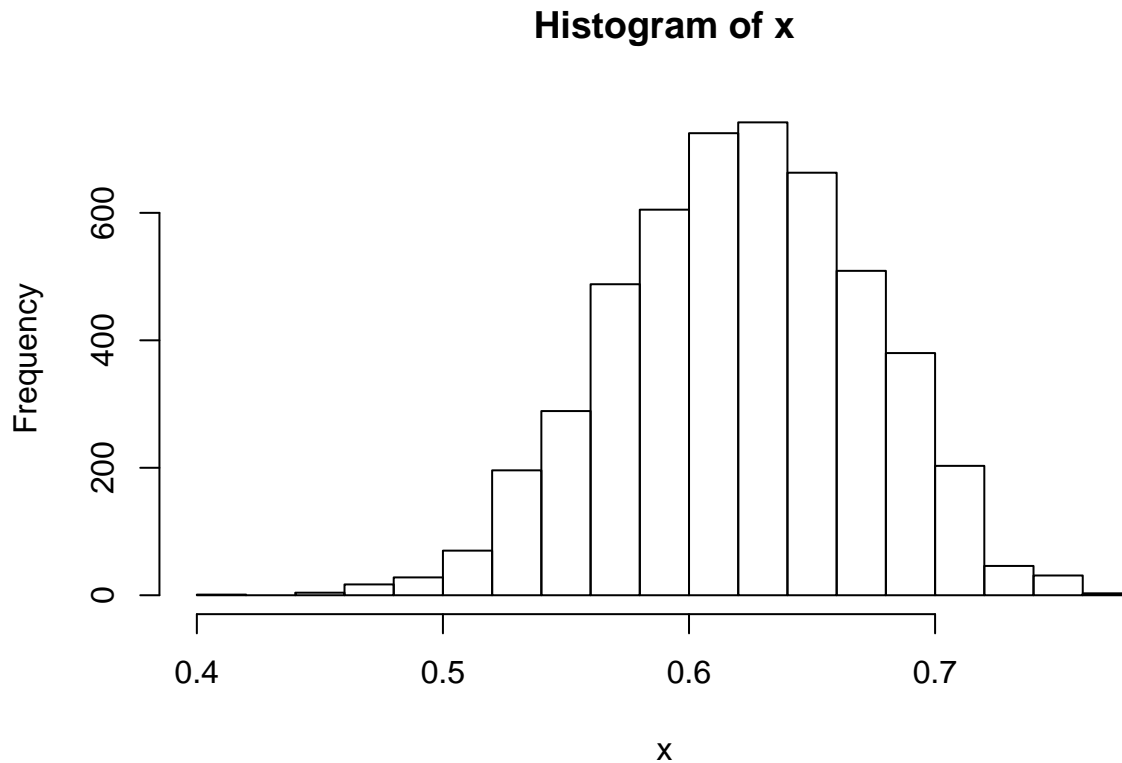
w <- 0.5 #width of uniform support set
m <- 5000 #length of the chain
burn <- 1000 #burn in time
animals <- 197

u <- runif(m)
v <- runif(m,-w,w)

x[1] <- .5
for(i in 2:m)
{
  y <- x[i-1] + v[i]
  if(u[i] <= prob(y,count)/prob(x[i-1],count))
    x[i] <- y else
    x[i] <- x[i-1]
}

hist(x)

```



```
xb <- x[(burn+1):m]
theta <- mean(xb)
```

```
theta
```

```
## [1] 0.6209864
```

```
#posterior distribution
```

```
p <- c(.5+theta/4, (1-theta)/4, (1-theta)/4, theta/4)
```

```
p
```

```
## [1] 0.65524659 0.09475341 0.09475341 0.15524659
```

```
#p*197
```

Hence the $\theta = 0.6187743$

Therefore the posterior distribution will be: (0.6552466, 0.0947534, 0.0947534, 0.1552466)

Problem 3:

λ = Mean of the data before the change point

ϕ = Mean of the data after the change point

m = Change point

β = Scale parameter for the distribution of λ

δ = Scale parameter for the distribution of ϕ

The following R code will obtain the following parameters (will run for 5000 iterations):

```
#Read the coal data set

library(boot)
data(coal)

year <- floor(coal)
y <- table(year)

#plot(y)

y <- floor(coal[[1]])
y <- tabulate(y)
y <- y[1851:length(y)]

#plot(y)

# Initialization
n <- length(y) #Length of data

m <- 5000 #length of chain

mu <- lambda <- k <- b1 <- b2 <- numeric(m)

L <- numeric(n)

k[1] <- sample(1:n,1)
mu[1] <- 1
lambda[1] <- 1
b1[1] <- 1
b2[1] <- 1

#Run the Gibbs sampler

for(i in 2:m)
{
  kt <- k[i-1]

  #generate mu
  r <- .5 + sum(y[1:kt])
  mu[i] <- rgamma(1, shape=r,rate=kt+b1[i-1])

  #generate lambda
  if(kt+1 > n)
    r <- .5 +sum(y)
```

```

else
  r <- .5 + sum(y[(kt+1):n])
  lambda[i] <- rgamma(1,shape=r,rate=n-kt+b2[i-1])

  #generate b1 and b2
  b1[i] <- rgamma(1,shape=.5,rate=mu[i]+1)
  b2[i] <- rgamma(1,shape=.5,rate=lambda[i]+1)

  for(j in 1:n)
  {
    L[j] <- exp((lambda[i] - mu[i]) * j) *
      (mu[i] / lambda[i])^sum(y[1:j])

  }

  L <- L / sum(L)

  #Generate k from discrete dist L on 1:n
  k[i] <- sample(1:n,prob=L, size=1)
}

phi <- lambda
lambda <- mu
m <- k
beta <- b1
delta <- b2

mean(phi)

```

```
## [1] 0.926059
```

```
sd(phi)
```

```
## [1] 0.1203184
```

```
mean(lambda)
```

```
## [1] 3.124374
```

```
sd(lambda)
```

```
## [1] 0.2945038
```

```
mean(m)
```

```
## [1] 39.9146
```

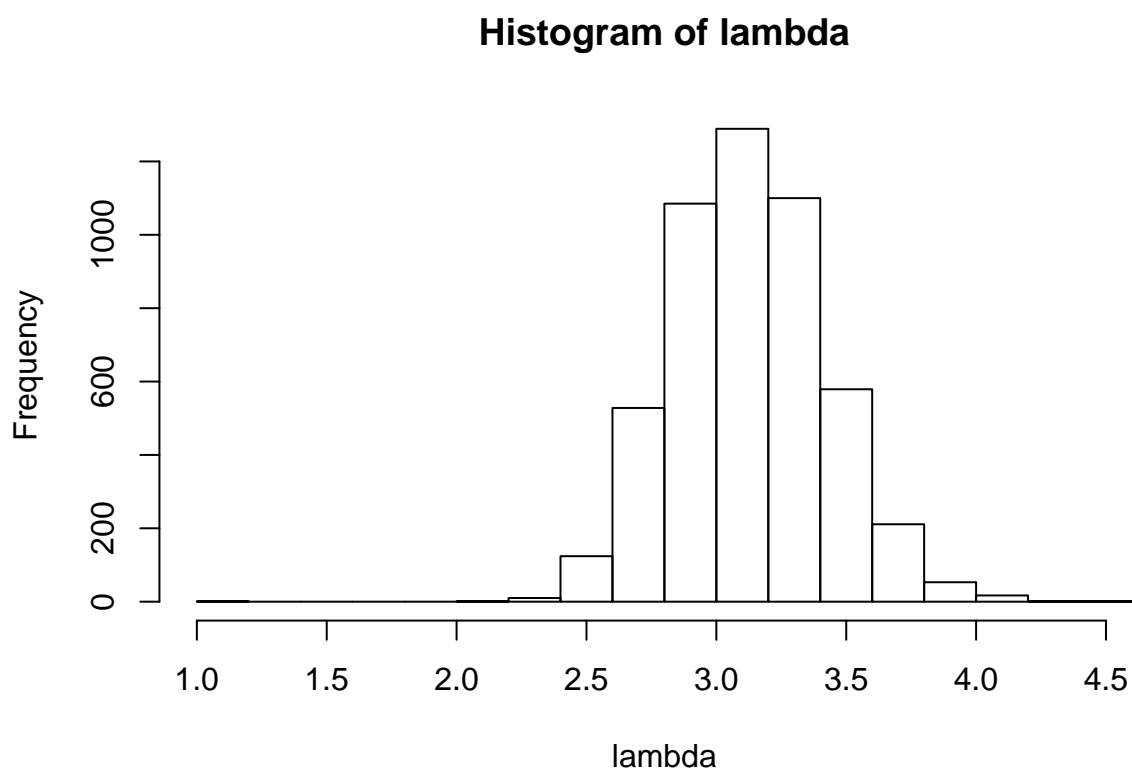


```
sd(m)
```

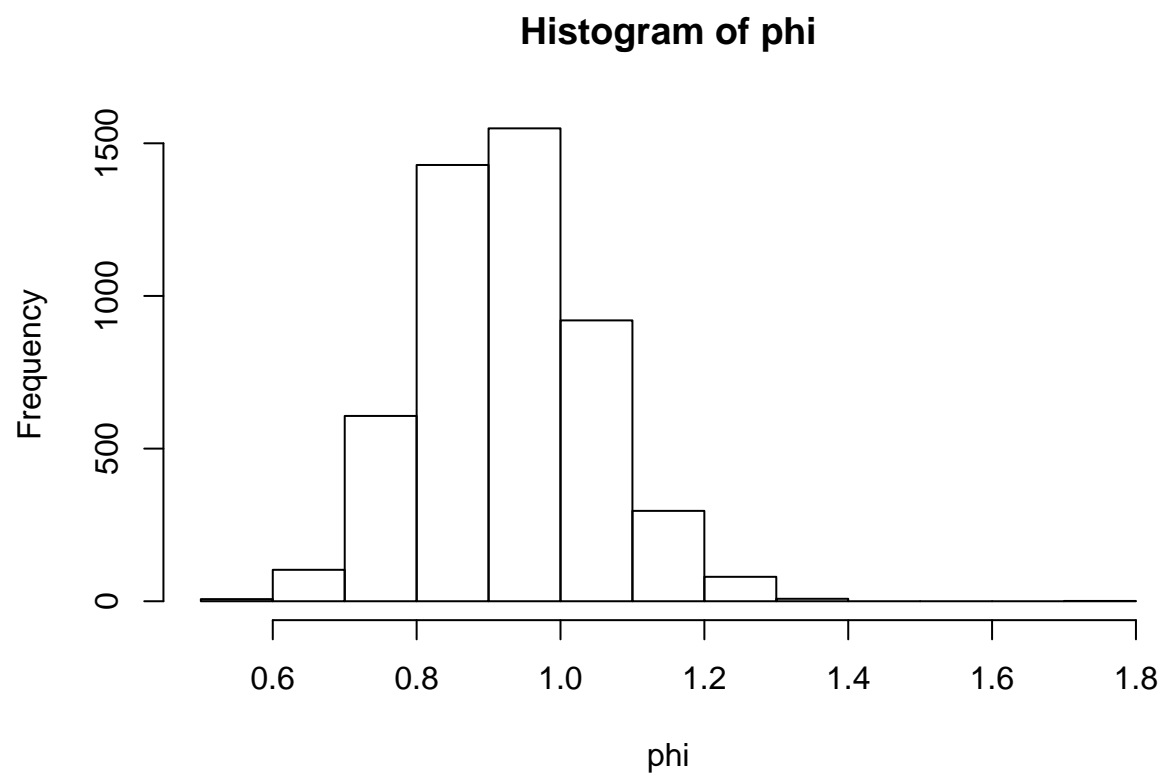
```
## [1] 2.504189
```

a)

```
#par(mfrow=c(1,1))  
#Histogram of lambda  
hist(lambda)
```

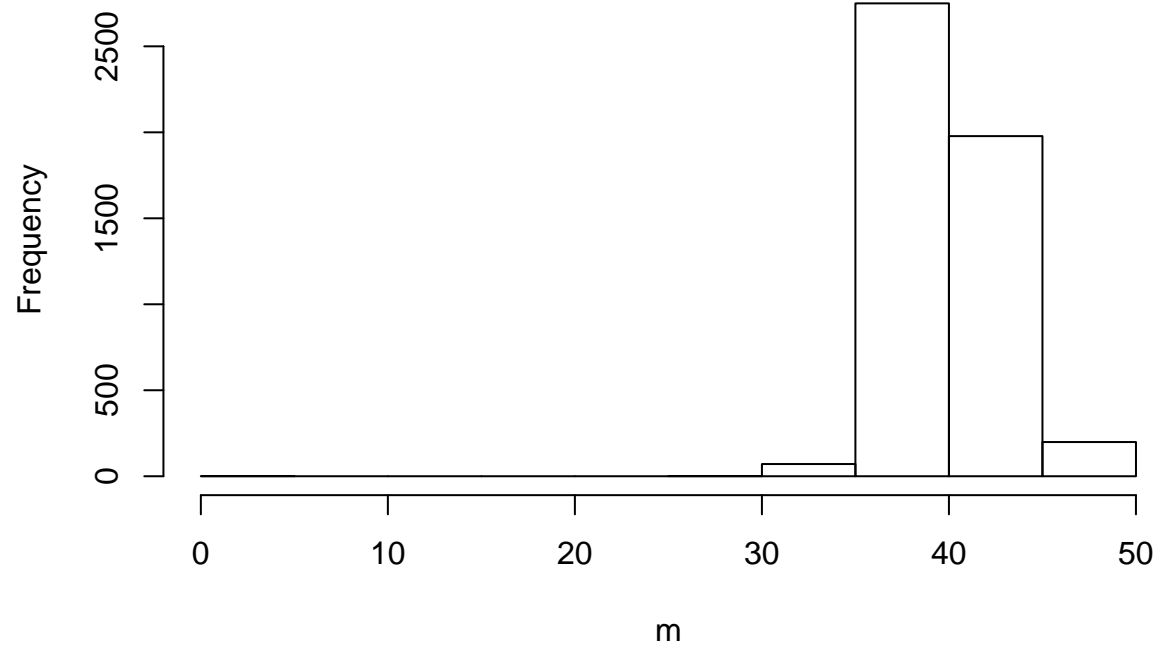


```
#Histogram of phi  
hist(phi)
```



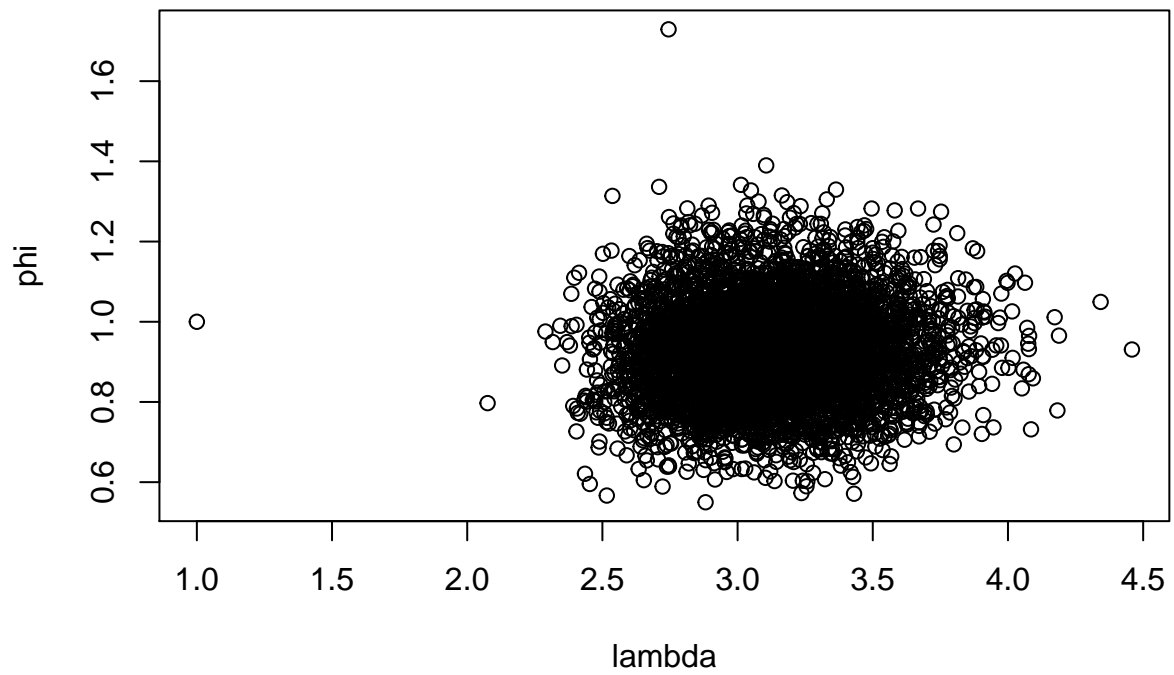
```
#Histogram of m  
hist(m)
```

Histogram of m



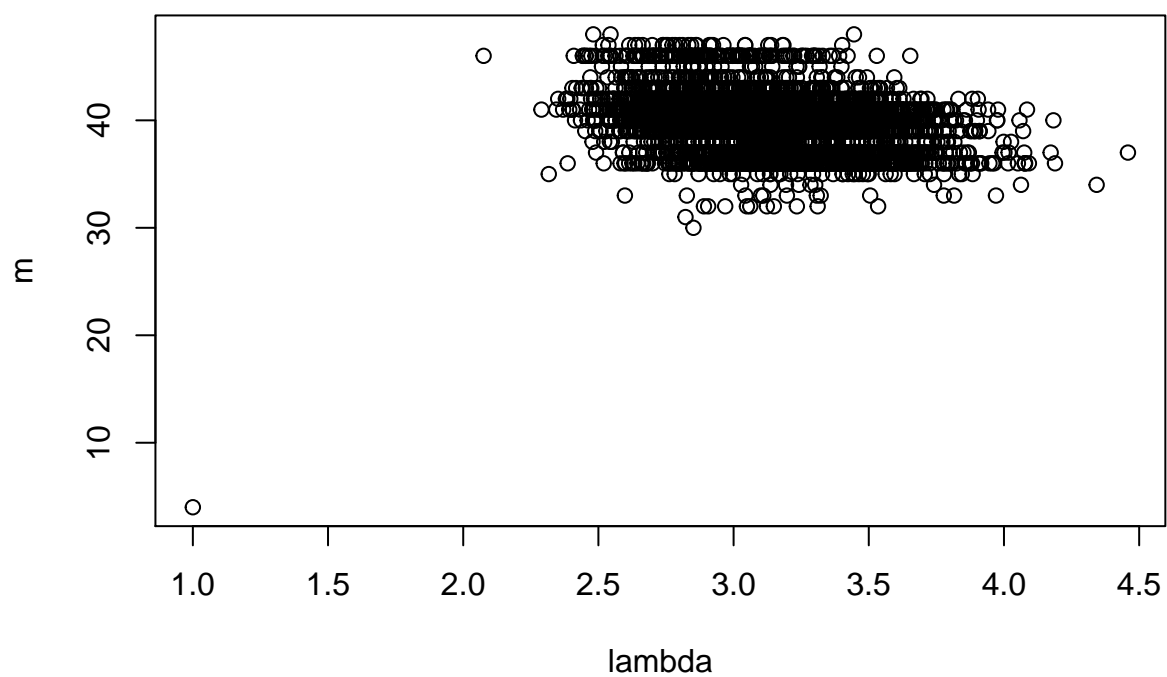
```
#Plot of lambda vs phi  
plot(lambda,phi,main="plot of lambda vs phi")
```

plot of lambda vs phi



```
#Plot of lambda vs m  
plot(lambda,m,main="plot of lambda vs m")
```

plot of lambda vs m



```
#Plot of beta vs delta  
plot(beta,delta,main="plot of beta vs delta")
```



b)

The change point has occurred after 39.9146 years. The 95% confidence interval is $39.9146 \pm (1.96) (2.5041887)$
 $= [35.0063902, 44.8228098]$

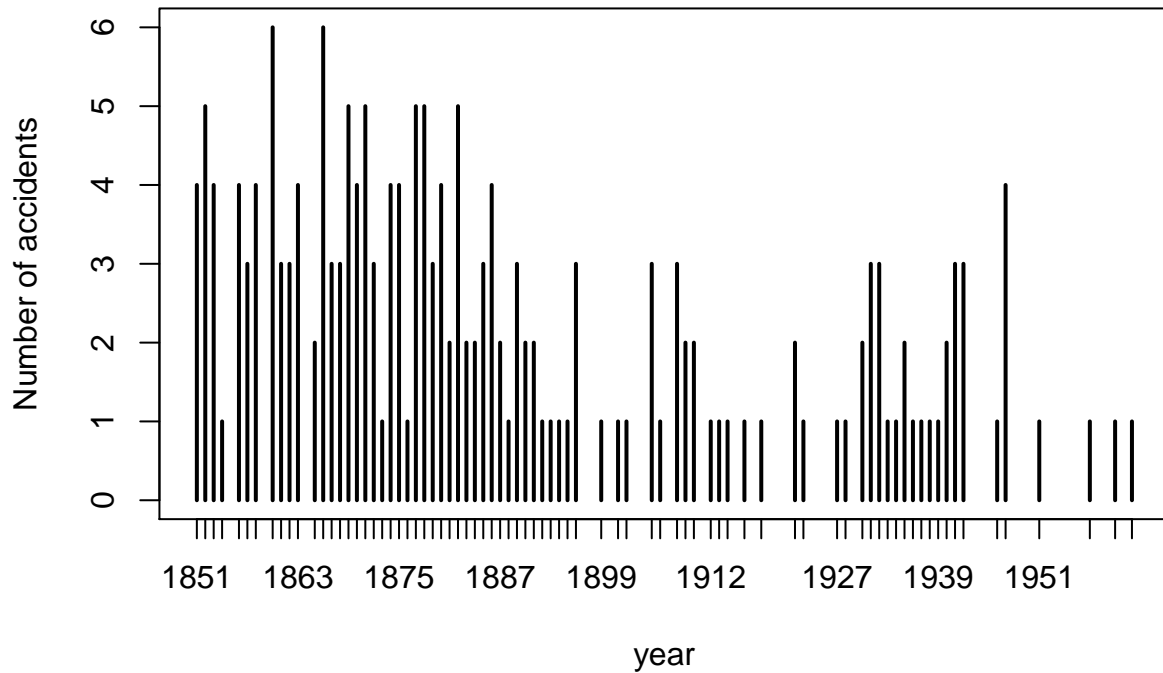
The average number of accidents before change point = 3.1243736/year and average number of accidents after change point = 0.926059/year

Let us plot the time series of the observations:

```
year <- floor(coal)
y <- table(year)

plot(y,main="Number of Accidents between 1851 and 1962",ylab="Number of accidents")
```

Number of Accidents between 1851 and 1962



We can clearly see that after approximately 40 years (around 1891), the number of accidents have drastically decreased. Our results obtained are consistent with the time series of observed data, since we see a drastic decrease in the accidents after 40 years.

c)

Gibbs sampling is a special case of Metropolis-Hastings method. In Gibbs sampling we do not reject any sample, and Gibbs sampling is often applied when the target distribution is multi-variate.

Problem-4

The cost matrix is displayed below:

cost													
##	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]
## [1,]	0	633	257	91	412	150	80	134	259	505	353	324	70
## [2,]	633	0	390	661	227	488	572	530	555	289	282	638	567
## [3,]	257	390	0	228	169	112	196	154	372	262	110	437	191
## [4,]	91	661	228	0	383	120	77	105	175	476	324	240	27
## [5,]	412	227	169	383	0	267	351	309	338	196	61	421	346
## [6,]	150	488	112	120	267	0	63	34	264	360	208	329	83
## [7,]	80	572	196	77	351	63	0	29	232	444	292	297	47
## [8,]	134	530	154	105	309	34	29	0	249	402	250	314	68

```
## [9,] 259 555 372 175 338 264 232 249 0 495 352 95 189
## [10,] 505 289 262 476 196 360 444 402 495 0 154 578 439
## [11,] 353 282 110 324 61 208 292 250 352 154 0 435 287
## [12,] 324 638 437 240 421 329 297 314 95 578 435 0 254
## [13,] 70 567 191 27 346 83 47 68 189 439 287 254 0
## [14,] 211 466 74 182 243 105 150 108 326 336 184 391 145
## [15,] 268 420 53 239 199 123 207 165 383 240 140 448 202
## [16,] 246 745 472 237 528 364 332 349 202 685 542 157 289
## [17,] 121 518 142 84 297 35 29 36 236 390 238 301 55
## [,14] [,15] [,16] [,17]
## [1,] 211 268 246 121
## [2,] 466 420 745 518
## [3,] 74 53 472 142
## [4,] 182 239 237 84
## [5,] 243 199 528 297
## [6,] 105 123 364 35
## [7,] 150 207 332 29
## [8,] 108 165 349 36
## [9,] 326 383 202 236
## [10,] 336 240 685 390
## [11,] 184 140 542 238
## [12,] 391 448 157 301
## [13,] 145 202 289 55
## [14,] 0 57 426 96
## [15,] 57 0 483 153
## [16,] 426 483 0 336
## [17,] 96 153 336 0
```

Let us write a function that finds the cost of a potential solution:

```
find_cost <- function(cost_matrix,solution)
{
  y <- numeric(16)
  for(i in 1:16)
  {
    y[i] <- cost_matrix[solution[i],solution[i+1]]
  }
  return(sum(y))
}
```

The simulated annealing algorithm is defined below:

```
simulated_annealing <- function(T0, drop,n,cost_matrix)
{
  #Get the number of cities
  s <- nrow(cost_matrix)

  #Declare a matrix to store all the solutions
  accepted_solution <- matrix(rep(0,170000),
                              byrow=TRUE,nrow=10000)
```



```

#Vector to store all costs
c_1 <- vector()

#Initial solution
solution_1 <- sample(1:s,s)

#Initial cost
c_1[1] <- find_cost(cost_matrix,solution_1)

#Initial accepted solution
accepted_solution[1,] <- solution_1

for(i in 2:n)
{
  #Get the 2 elements randomly
  index <- sort(sample(1:s,s)[1:2])

  #Prepare a second solution
  solution_2 <-
    c(solution_1[1:ifelse(index[1] == 1, 1,
                          (index[1]-1))],
      solution_1[(ifelse(index[2]==17,16,
                          index[2])):ifelse(index[1]==1,2,index[1])],
      solution_1[ifelse((index[2])==17,17,
                          (index[2]+1)):17])

  #Find the cost of the current solution or solution-2
  c_1[i] <- find_cost(cost_matrix,solution_2)

  #If solution-2 is better than solution-1, then accept solution-2 blindly
  if(c_1[i-1] > c_1[i])
  {
    accepted_solution[i,] <- solution_2
    T0 <- T0 * drop
    solution_1 <- solution_2
  } else
  {
    p <- exp((c_1[i-1] - c_1[i])/T0)
    u <- runif(1)
    if(p > u)
    {
      accepted_solution[i,] <- solution_2
    }
    else{
      c_1[i] <- c_1[i-1]
      accepted_solution[i,] <- solution_1
    }
    T0 <- T0 * drop
  }
}

return(list(solution=accepted_solution,cost=c_1,T0=T0))

```

```

}

#Get the cost of the best solution
l <- simulated_annealing(1, 0.9999,10000,cost)

l$cost[10000]

```

```
## [1] 1948
```

```

#Let us get the cost of a random path
x <- sample(1:17,17)

find_cost(cost,x)

```

```
## [1] 3307
```

The cost of the best solution is 1948, while the cost of the random path is: 3307.

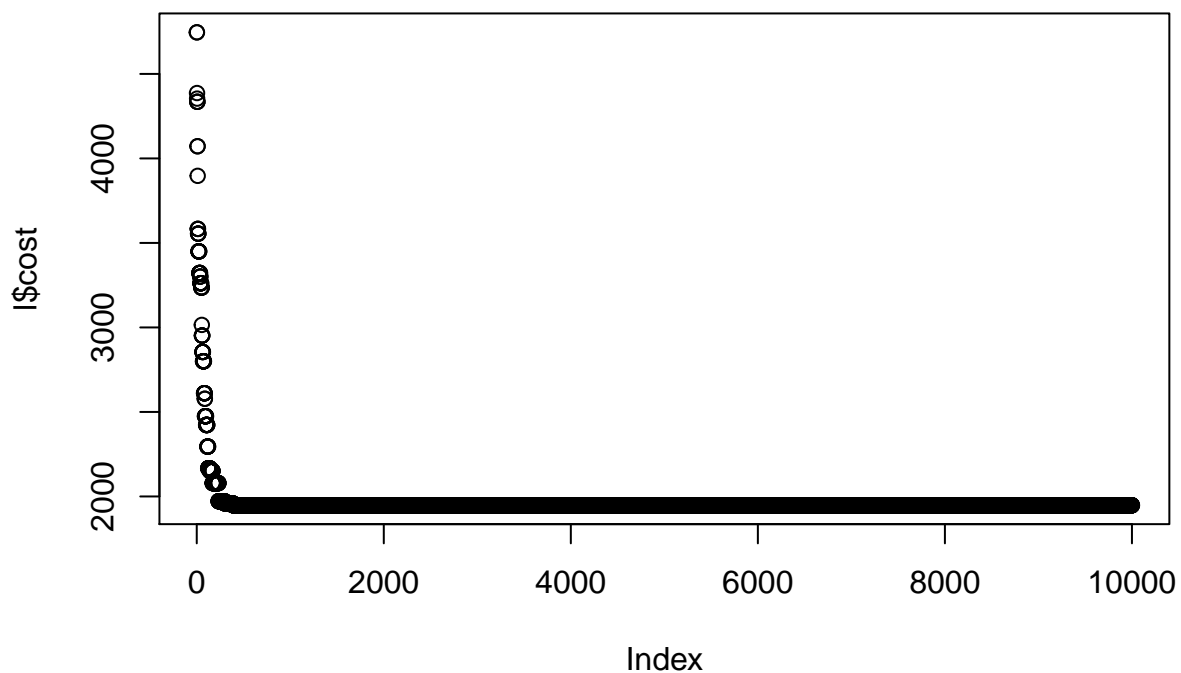
The optimal path obtained using the algorithm is:

```
l$solution[10000,]
```

```
## [1] 11 5 2 10 15 3 14 6 17 13 4 9 12 16 1 7 8
```

Let us plot the iterations vs the cost of the solution in the respective iteration:

```
plot(l$cost)
```



We can infer that after approximately 300 iterations, the cost has stabilized.

Repeating the simulated annealing for 4 times to find how the cost of the optimal solution is varying

```
set.seed(1234)
l1 <- simulated_annealing(1, 0.9999, 10000, cost)
l2 <- simulated_annealing(1, 0.9999, 10000, cost)
l3 <- simulated_annealing(1, 0.9999, 10000, cost)
l4 <- simulated_annealing(1, 0.9999, 10000, cost)

l1$solution[10000,]

## [1] 2 10 5 11 3 6 8 7 1 16 12 9 4 13 17 14 15

l1$cost[10000]

## [1] 1819

l2$solution[10000,]

## [1] 3 6 8 7 1 13 4 16 12 9 17 14 15 11 10 2 5

l2$cost[10000]

## [1] 2040

l3$solution[10000,]

## [1] 15 3 14 17 6 8 7 13 4 1 16 12 9 11 5 2 10

l3$cost[10000]

## [1] 1913

l4$solution[10000,]

## [1] 9 12 16 4 13 17 14 15 10 2 5 11 3 6 8 7 1

l4$cost[10000]

## [1] 1906
```

The optimal solution's cost does not vary a lot, but the optimal solution is drastically different in different repetitions.